

kNN Clustering - Classification

Karen Mazidi

This example shows how to do knn clustering for classification.

The iris database comes with R. It has 150 instances and 5 columns: - Sepal.Length - Sepal.Width - Petal.Length - Petal.Width - Species: setosa, versicolor or virginica

Load and look at the data

```
attach(iris)
str(iris)      # display the structure of the object

## 'data.frame':   150 obs. of  5 variables:
## $ Sepal.Length: num  5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
## $ Sepal.Width : num  3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
## $ Petal.Length: num  1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
## $ Petal.Width : num  0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
## $ Species      : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 1 ...

summary(iris)


##      Sepal.Length      Sepal.Width      Petal.Length      Petal.Width
## Min.      :4.300    Min.      :2.000    Min.      :1.000    Min.      :0.100
## 1st Qu.:5.100    1st Qu.:2.800    1st Qu.:1.600    1st Qu.:0.300
## Median :5.800    Median :3.000    Median :4.350    Median :1.300
## Mean   :5.843    Mean   :3.057    Mean   :3.758    Mean   :1.199
## 3rd Qu.:6.400    3rd Qu.:3.300    3rd Qu.:5.100    3rd Qu.:1.800
## Max.   :7.900    Max.   :4.400    Max.   :6.900    Max.   :2.500
##      Species
## setosa      :50
## versicolor:50
## virginica  :50
##
##
##
```

Plot the data

We let the 3 classes show as 3 different colors with the bg parameter and the “unclass” values 1, 2, 3 representing the 3 types of irises.

```
plot(Petal.Length, Petal.Width, pch=21, bg=c("red", "green3", "blue")
      [unclass(Species)], main="Iris Data")
```

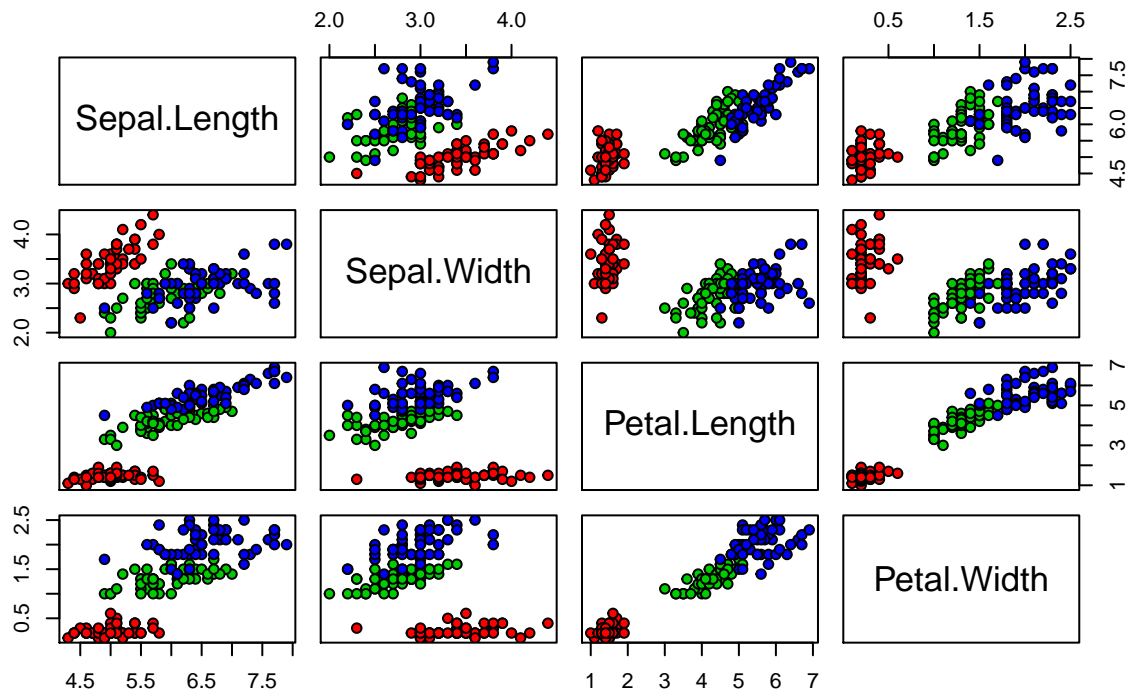
Iris Data



Pairs scatter plots

```
pairs(iris[1:4], main = "Iris Data", pch = 21, bg = c("red", "green3", "blue")[unclass(Species)])
```

Iris Data



Divide into train/test sets

We will randomly sample the data set to let 2/3 be training and 1/3 test,

```
set.seed(1958) # setting a seed gets the same results every time
ind <- sample(2, nrow(iris), replace=TRUE, prob=c(0.67, 0.33))
iris.train <- iris[ind==1, 1:4]
iris.test <- iris[ind==2, 1:4]
iris.trainLabels <- iris[ind==1, 5]
iris.testLabels <- iris[ind==2, 5]
```

Classify

The knn() function uses Euclidean distance to find the k nearest neighbors.

Classification is decided by majority vote with ties broken at random.

Using an odd k can avoid some ties.

```
library(class)
iris_pred <- knn(train=iris.train, test=iris.test, cl=iris.trainLabels, k=3)
```

Compute accuracy

We built a classifier with 98% accuracy.

It's often a good idea to scale the variables for clustering to make the distance calculations better. However in this case, the 3 predictors are roughly in the same scale so it's probably not necessary.

```
results <- iris_pred == iris.testLabels
acc <- length(which(results==TRUE)) / length(results)
# or combine into one line:
#acc <- length(which(iris_pred == iris.testLabels)) / length(iris_pred)
acc
```

```
## [1] 0.98
```